

HDFS

Senthil Kumar A

DataDotZ




Basics — *optional*

- Program
 - sequence of instructions written to perform a specified task with a computer
 - or a piece of code
- Process
 - an instance of a computer program that is being executed.
 - or a execution of a program
- Daemon Process
 - process which runs in background and has no controlling terminal.
- JVM – Java Virtual Machine
 - program which executes certain programs, namely those containing Java bytecode instructions

Basics — *optional*

- Client-server Concept

- Client sends requests to one or more servers which in turn accepts, processes them and return the requested information to the client.
- A server might run a software/program which listens on particular ip and port number for requests
 - Servers can also be said as daemons
- Examples:
 - Server - web server
 - Client – web browser

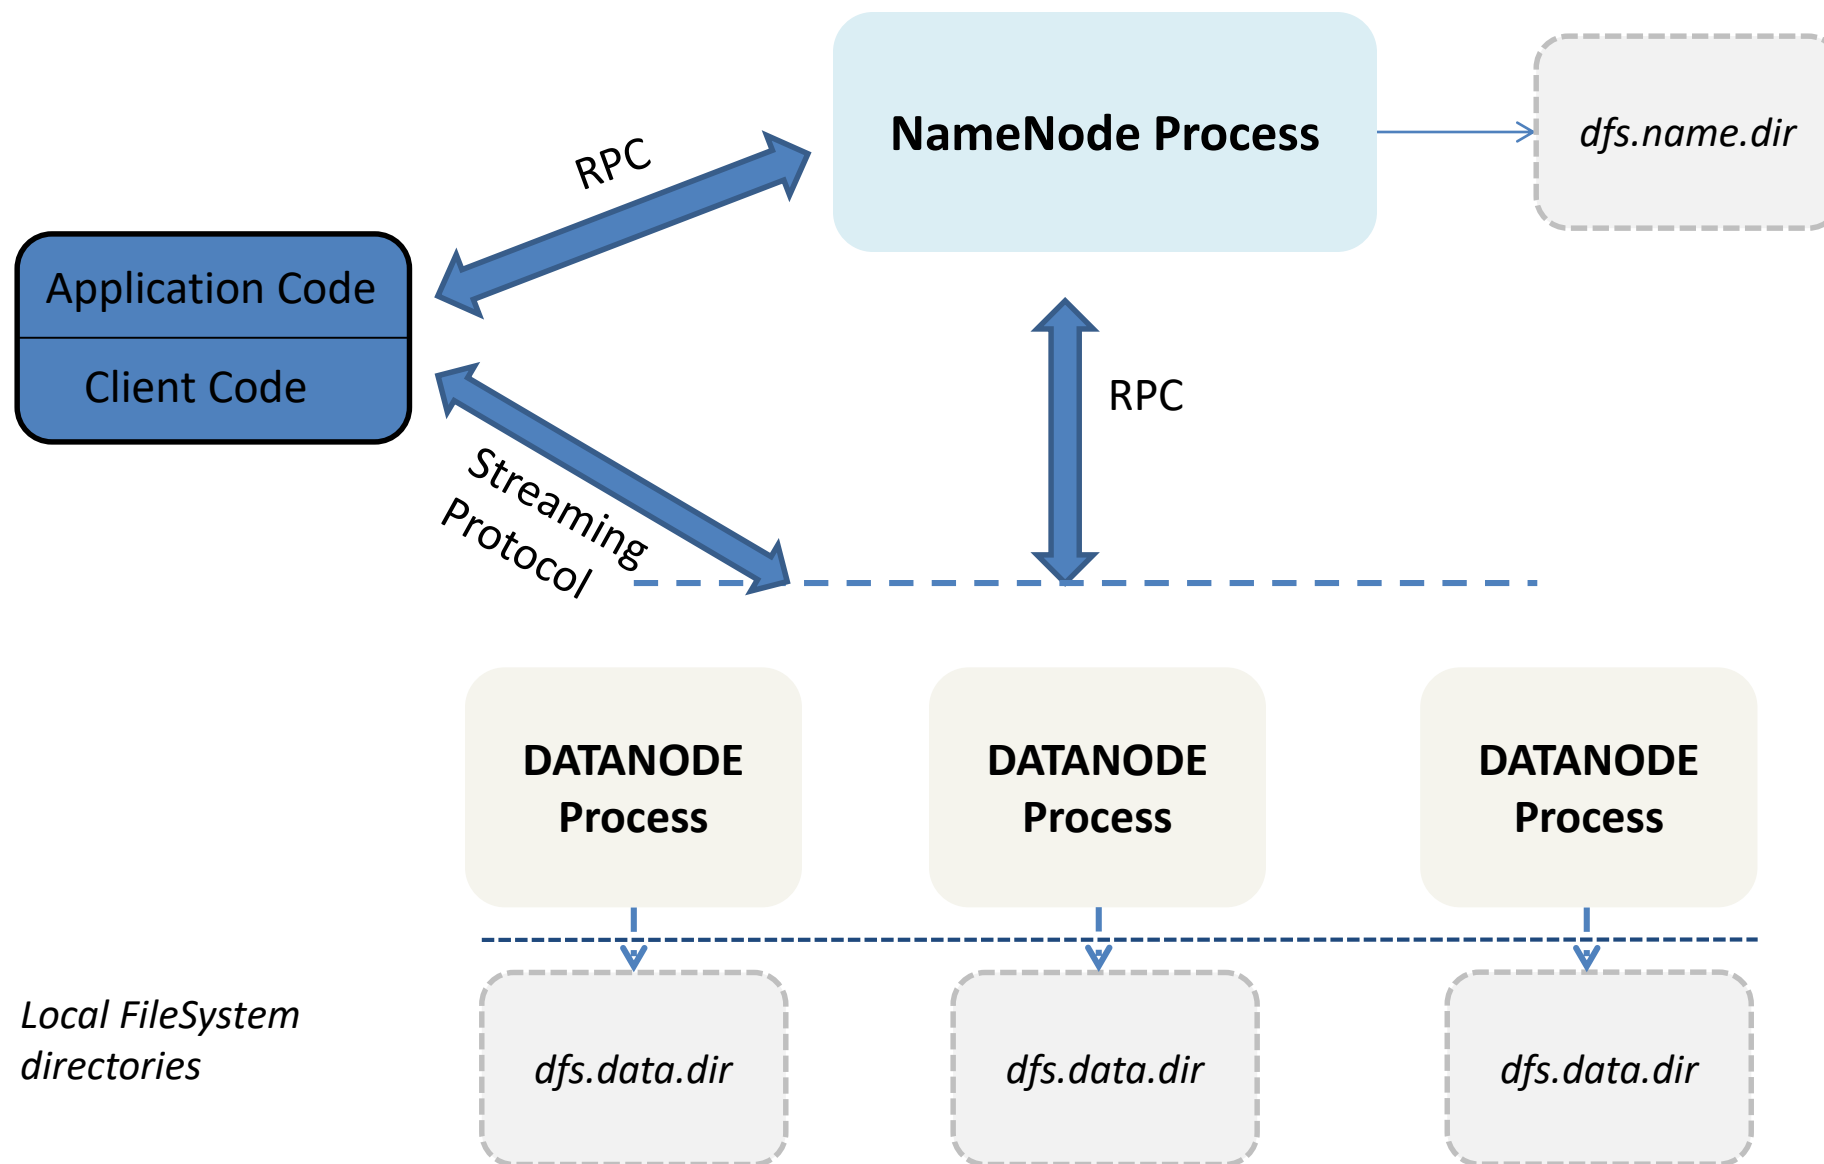


Introduction

- A distributed File System - **STORAGE**
 - A File System on multiple machines which sits on native filesystem
 - ext4,ext3
- Hardware Failure
 - Due to usage of Commodity machines, failure is a common phenomenon
 - Designed for failure
- Large Data Sets
 - Small Files Problem Due to NameNode
- Simple Coherency Model
 - Write Once , Read Many Times
- Streaming Data Access
 - High Throughput instead of low latency access


- ext3? ext4?

Continued...



Daemons in Hadoop Core

- NameNode
- DataNode
- Secondary NameNode
- Resource Manager*
- Node Manager*



* - will be seen later during MR

Block Concept

TestFile1.txt -> 1GB
Block Size -> 64 MB

Files are splitted into number of
chunks(Blocks) of pre-defined size

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B1,B2,.....B16

DataNode

B1
B8
B10
B13

DataNode

B3
B7
B12
B16

DataNode

B4
B5
B11
B14

DataNode

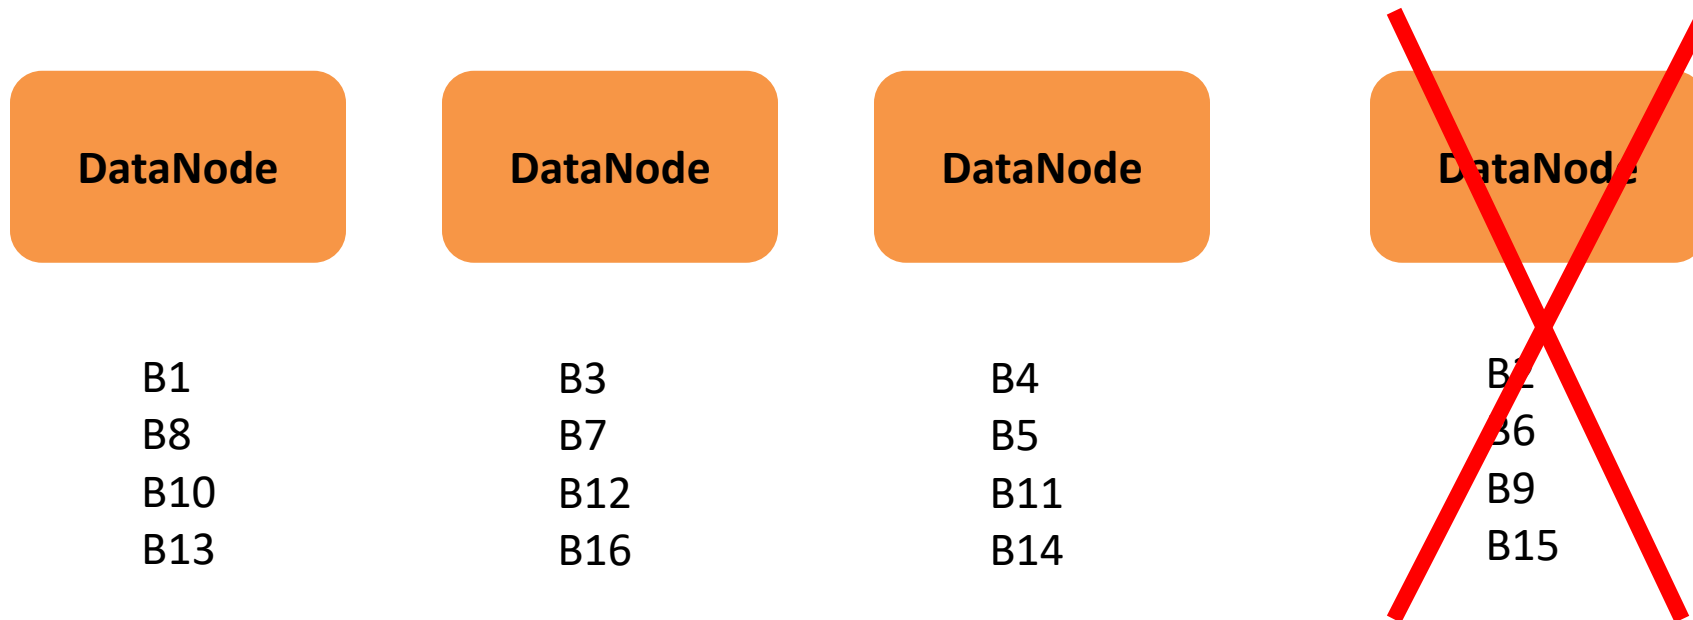
B2
B6
B9
B15

Block Concept

TestFile1.txt -> 1GB
Block Size -> 64 MB

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B1,B2,.....B16

**What happens to my data if
node 4 goes down??**

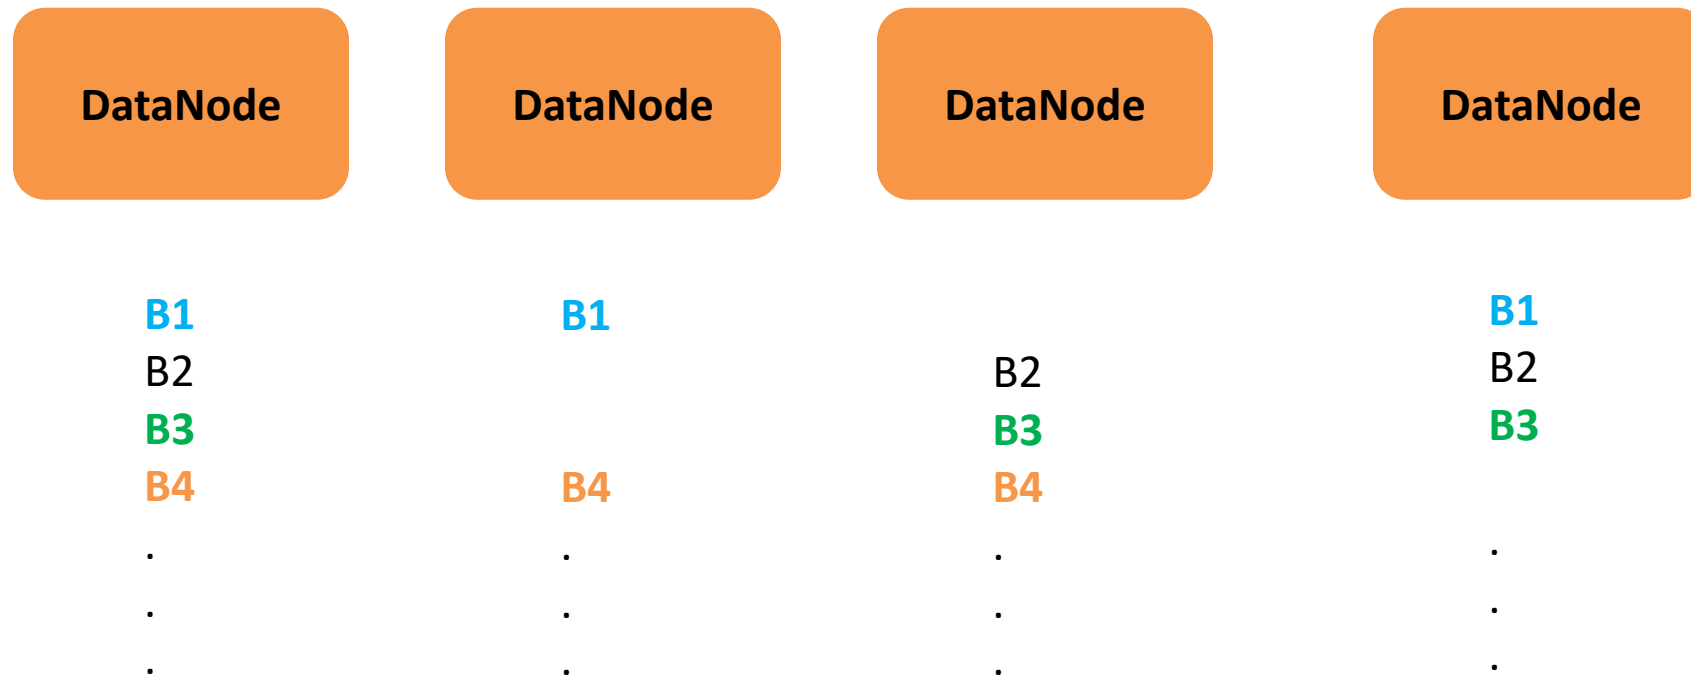


Fault Tolerant in HDFS

TestFile1.txt -> 1GB
Block Size -> 64 MB

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B1,B2,.....B16

HDFS provides fault tolerant by
replication of each block by 3

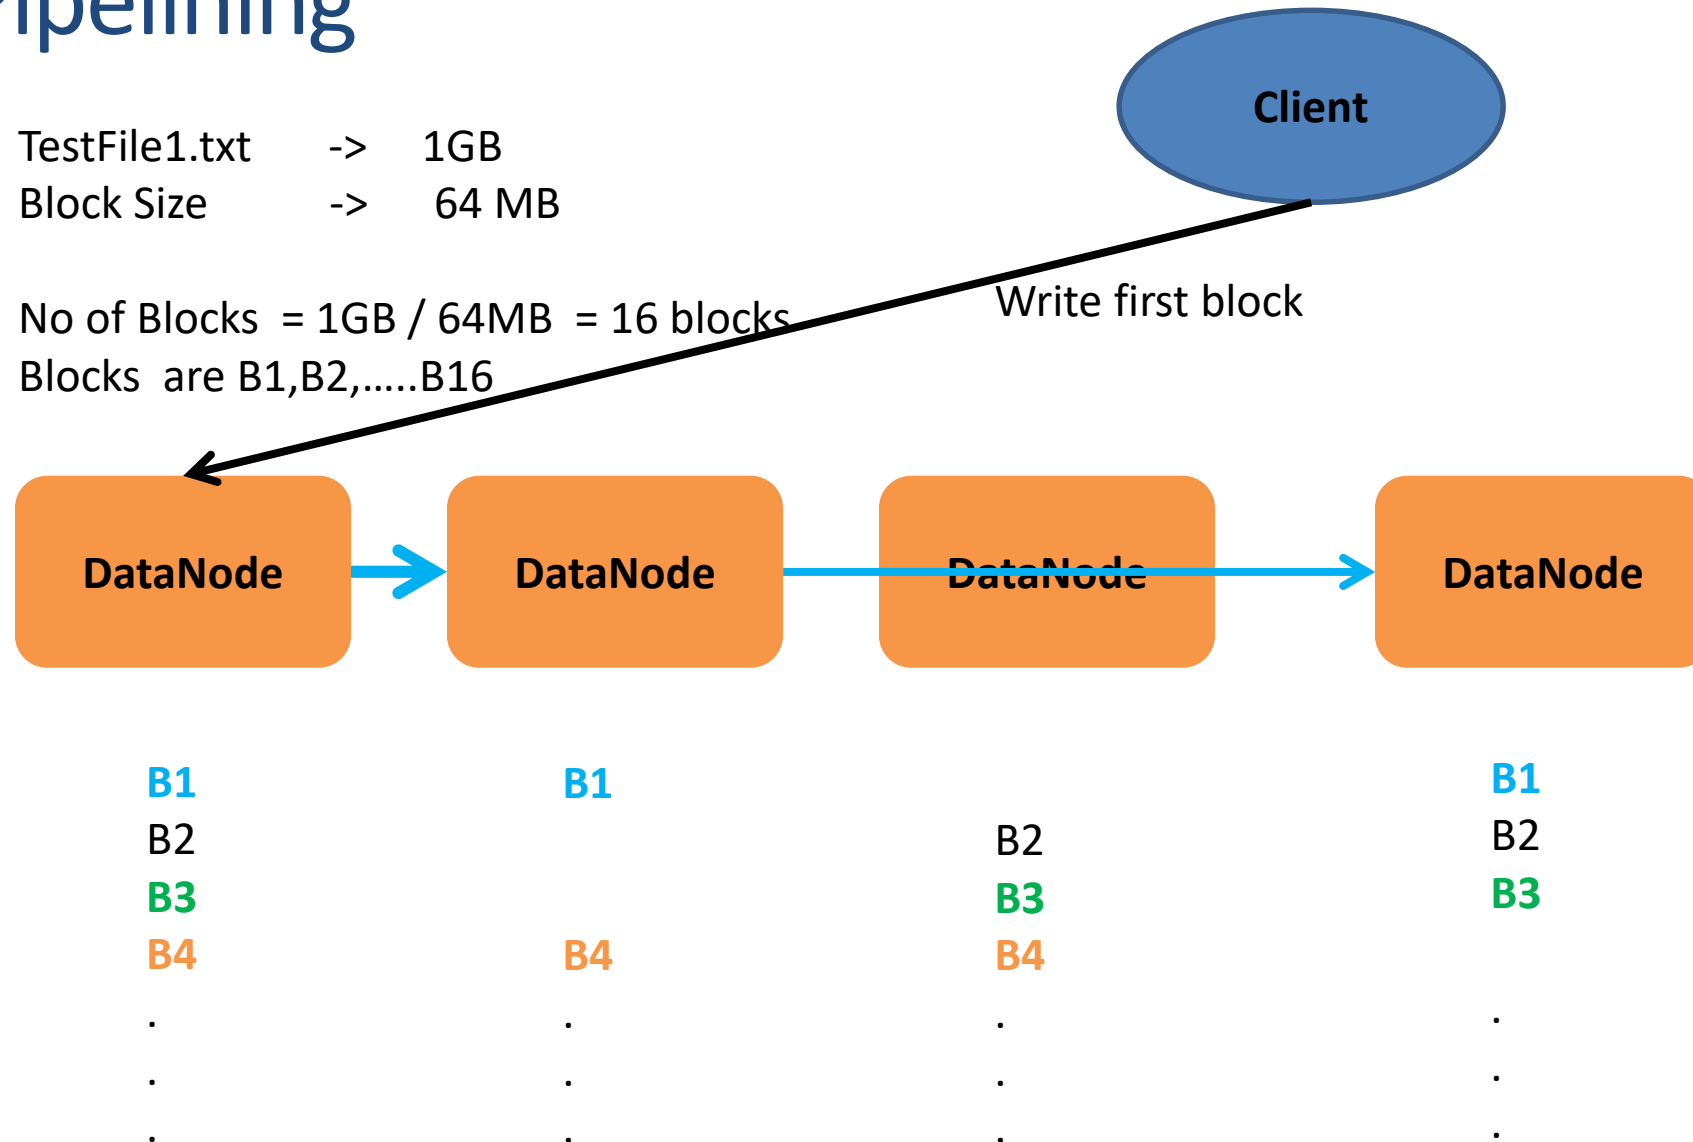


Data Pipelining

TestFile1.txt -> 1GB
Block Size -> 64 MB

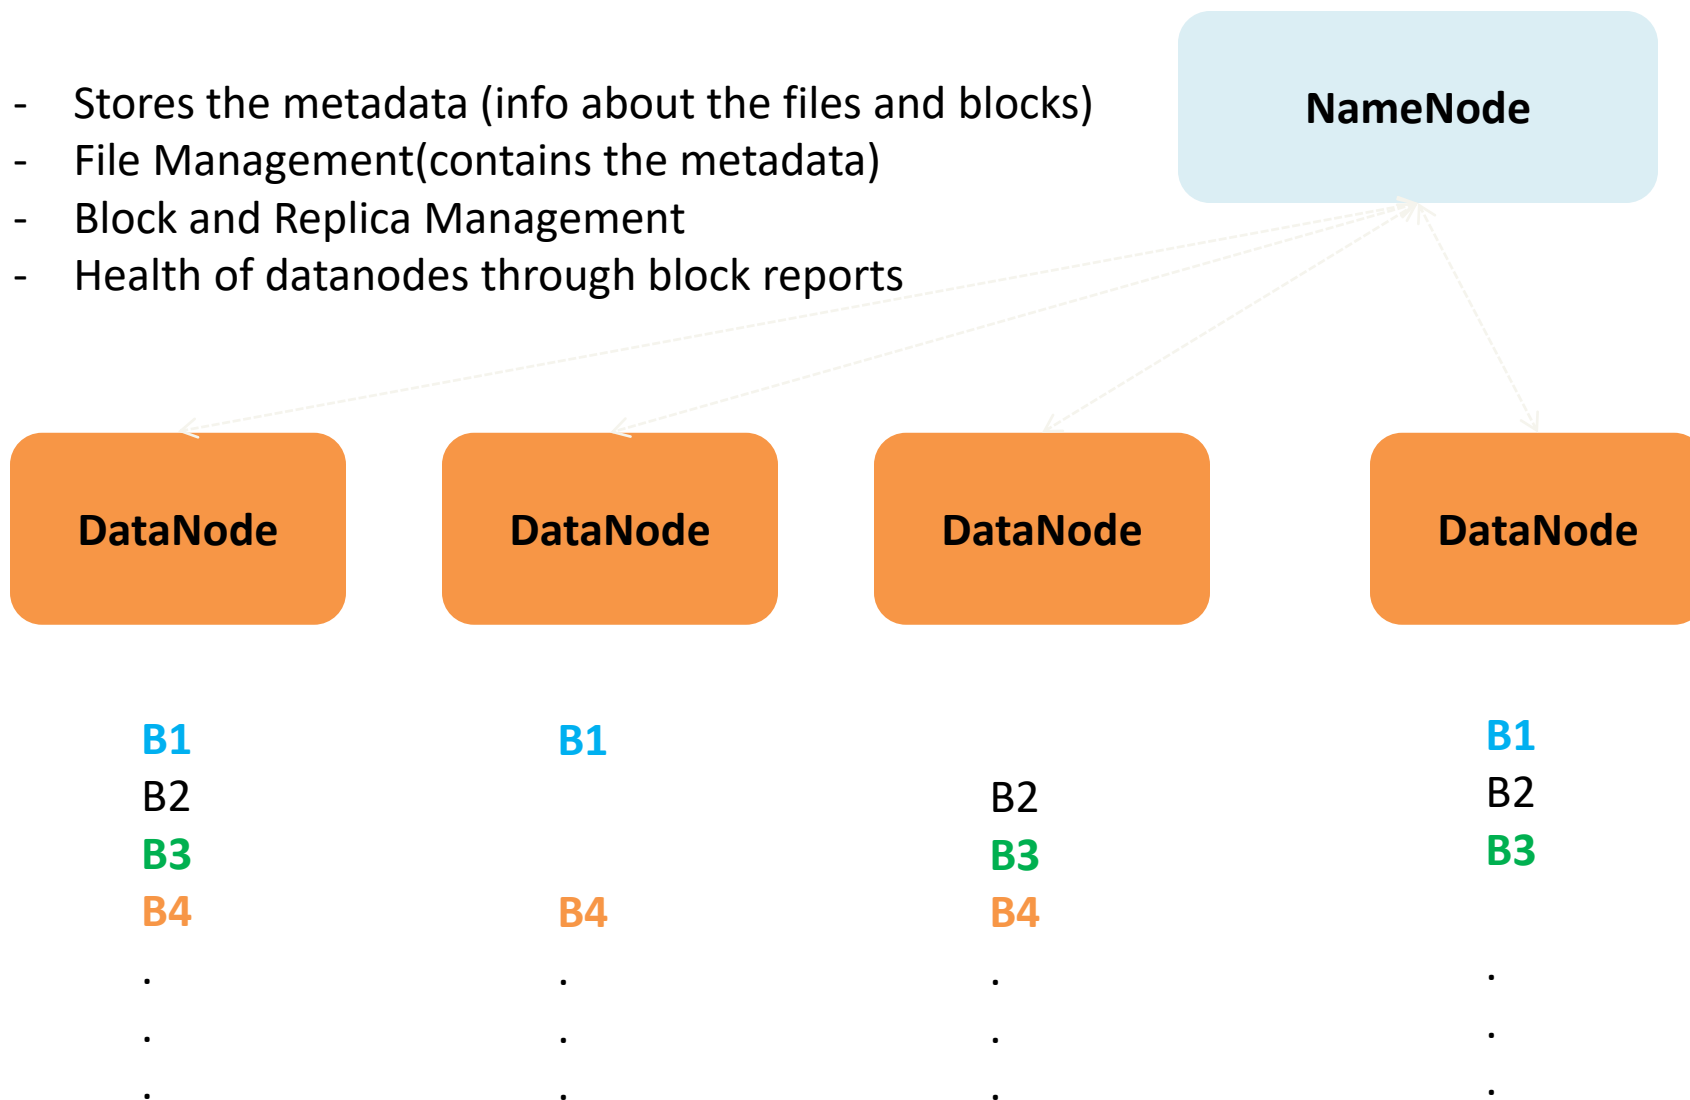
No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B1,B2,.....B16

Write first block

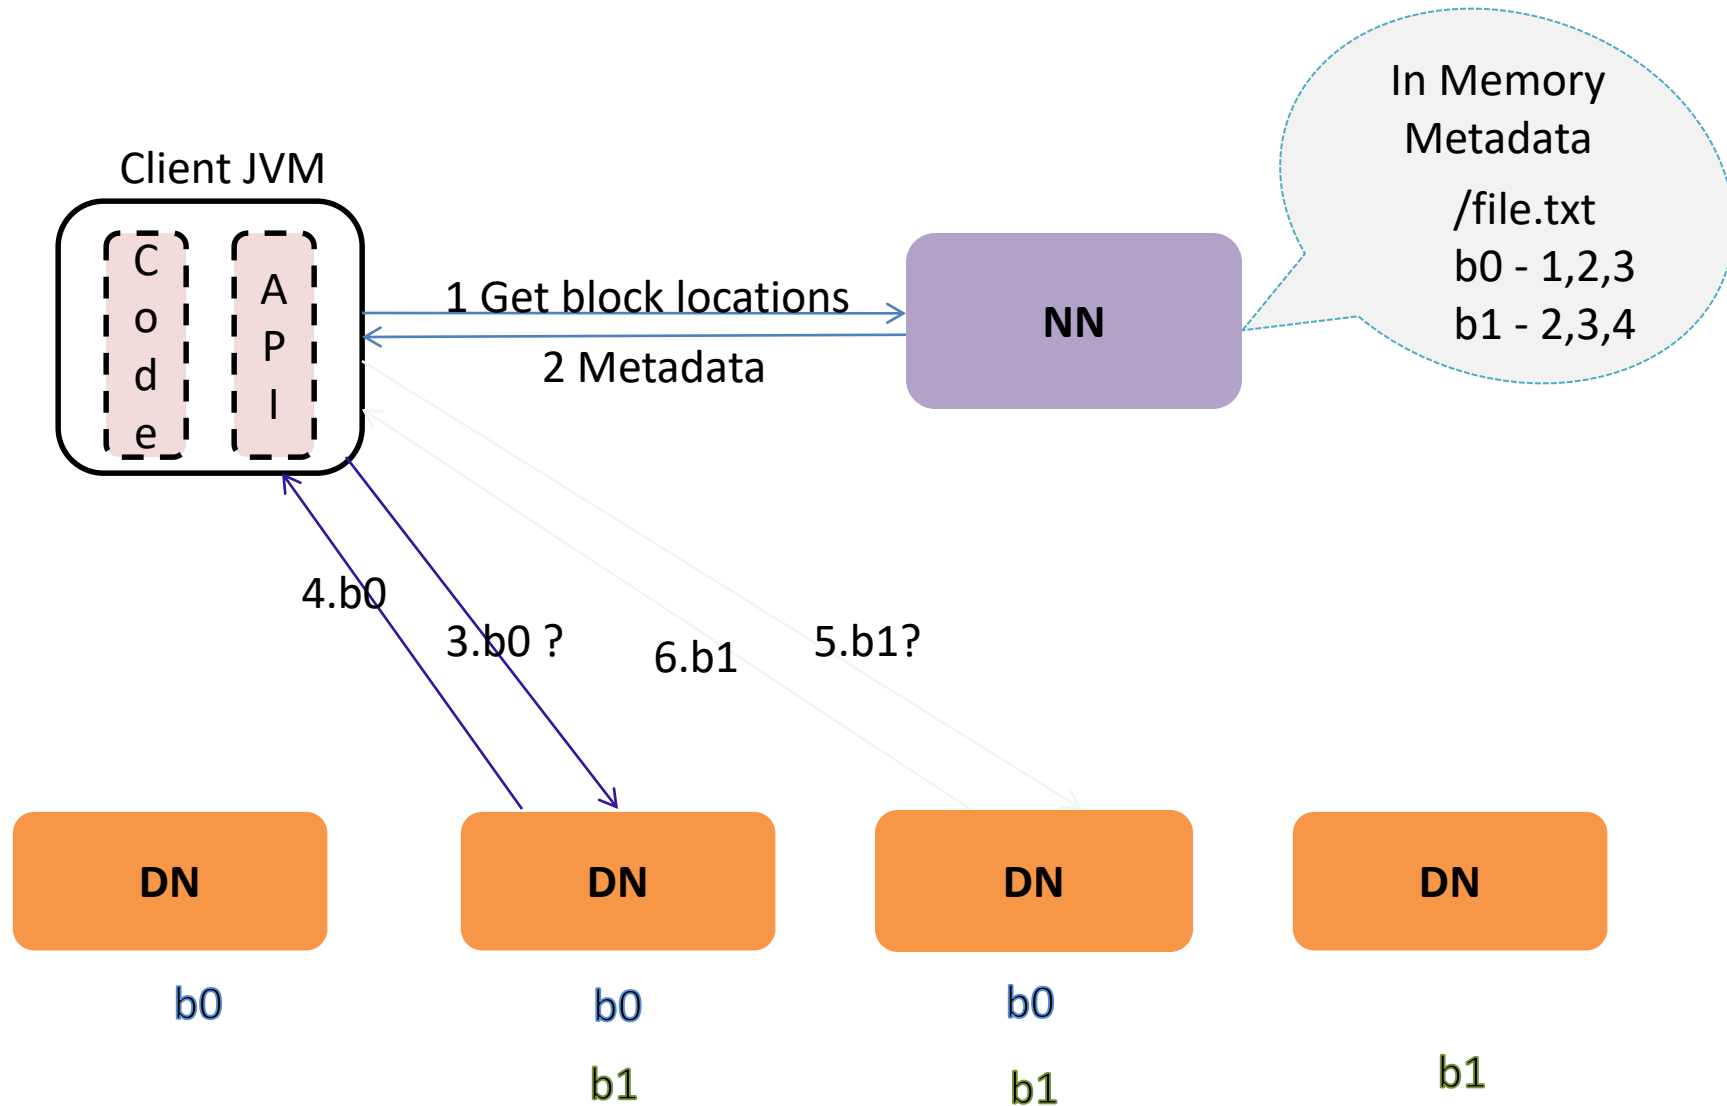


Role of NameNode

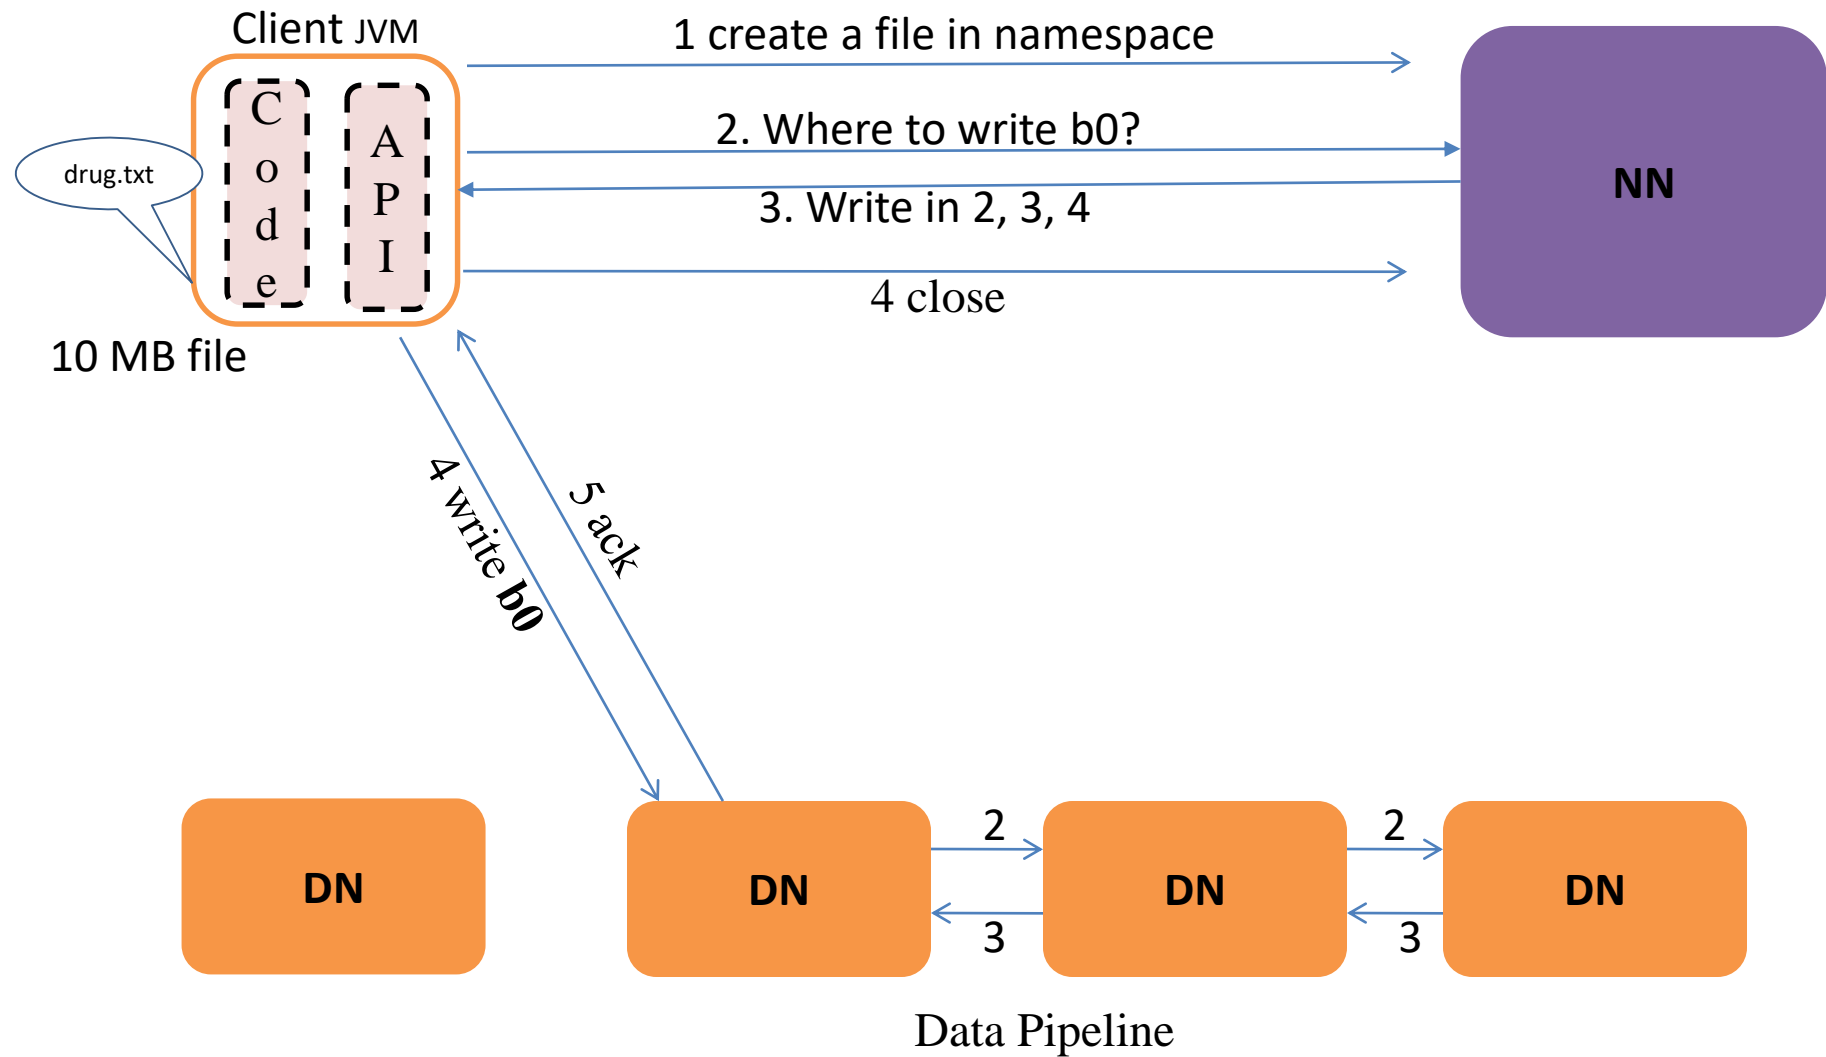
- Stores the metadata (info about the files and blocks)
- File Management(contains the metadata)
- Block and Replica Management
- Health of datanodes through block reports



File Read - Flow



File Write - Flow




Coherency in HDFS

- Write once, Read many times
- Completed blocks will be visible to users.
- Current block being in written will not be visible to users

Execution Modes & Installation


- Modes
 - Single Stand Alone
 - All Process runs in a single jvm
 - Does not use HDFS
 - Pseudo Distributed Mode - **for our training**
 - All daemon process runs in separate jvm in a single local machine
 - Used for development and testing
 - Uses HDFS to store data
 - Distributed Mode
 - A cluster of nodes more than 1
 - Each Process may run in different nodes
- Please follow instructor and doc provided

Minimal Configuration (pseudo distributed cluster)




SINGE NODE CLUSTER

Pseudo Distributed For Learning purpose




3 NODES CLUSTER



5 NODES CLUSTER



20 NODES + CLUSTERS



Ports used by Hadoop Daemons

Daemon	RPC	WEB
NameNode	8020 (50000*)	50070
SecondaryNameNode		50090
Resourcemanager	8032	8088
Node Manager	8040 , 8041	8042
DataNode	50010	50075
Mapreduce Job History Server		19888

After installation

- jps
 - JPS - JVM profiling status tool
- Web UI
 - NameNode - <http://localhost:50070>
 - Resource Manager – <http://localhost:8088>

Accessing HDFS

- Command line
 - Usage: *hadoop dfs <command>*
- JAVA API
- webHDFS

HDFS commands

- *hadoop dfs -copyFromLocal <srcLOCALfile> <destHDFSfile>*
- *hadoop dfs -ls /*
- *hadoop dfs -cat /<destHDFSfile>*
- *hadoop dfs -copyToLocal <srcHDFSfile> <destLOCALfile>*
- *hadoop dfs -mkdir /test*
- *hadoop dfs -rmr /test*

- *Please follow the document given*

JAVA API

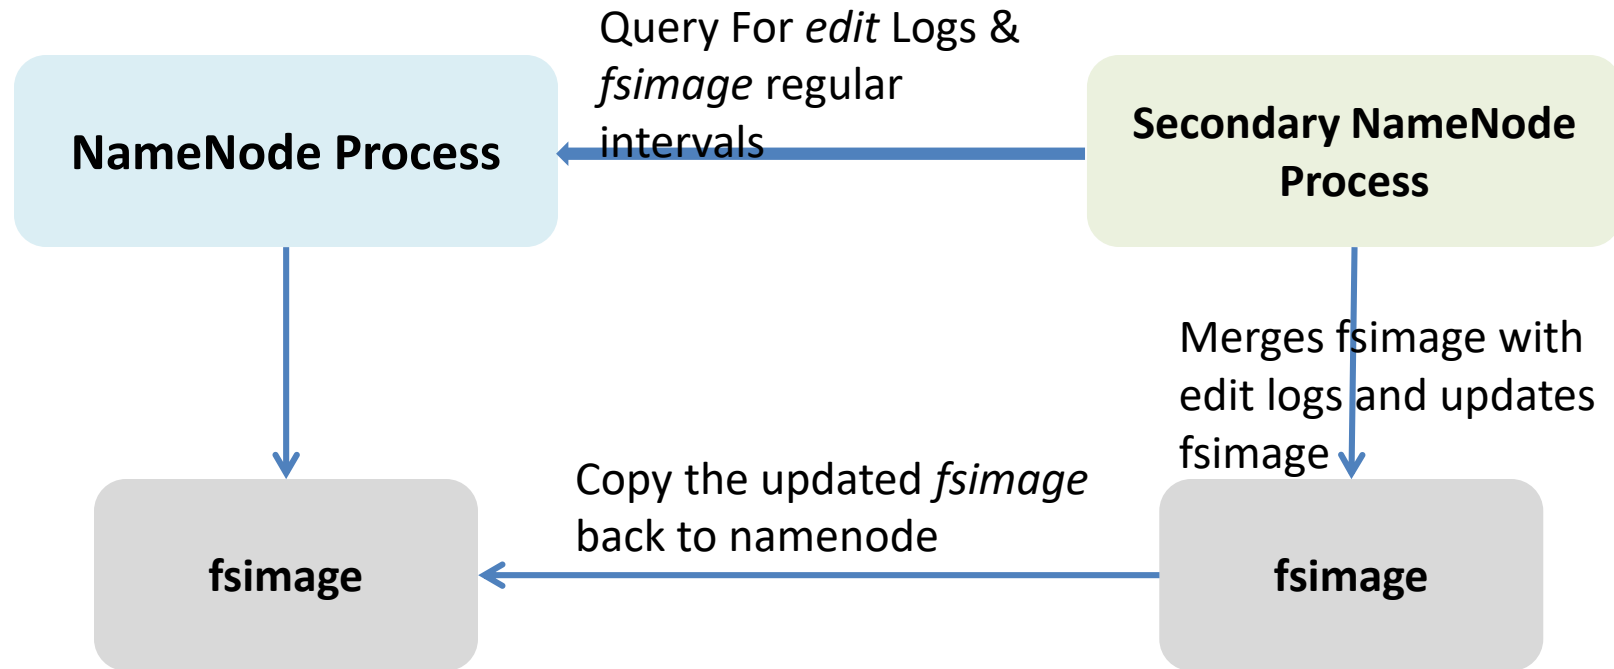
- Most Packages Used
 - *org.apache.hadoop.conf.Configuration*
 - *org.apache.hadoop.fs.BlockLocation*
 - *org.apache.hadoop.fs.FSDataInputStream*
 - *org.apache.hadoop.fs.FSDataOutputStream*
 - *org.apache.hadoop.fs.FileStatus*
 - *org.apache.hadoop.fs.FileSystem*
 - *org.apache.hadoop.fs.Path*
 - *org.apache.hadoop.hdfs.DistributedFileSystem*
 - *org.apache.hadoop.hdfs.protocol.DatanodeInfo*

Please see and execute the example code provided

FileSystem API methods

- `append()`
- `copyFromLocalFile()`
- `create()`
- `delete()`
- `makedirs()`
- `open()`

Role of Secondary NameNode



Rack AwareNess Policy

